

**Reference Manual
for the *rv32_cpu* RISC-V
Verilog softcore CPU
(version 0.2 draft)**



Simon Southwell

September 2021

Copyright

Copyright © 2021 Simon Southwell (Wyvern Semiconductors)


This document may not, in whole or part, be copied, photocopied, reproduced, translated, or reduced to any electronic medium or machine-readable form without prior written consent from the copyright holder.

Disclaimers

No warranties: the information provided in this document is “as is” without any express or implied warranty of any kind including warranties of accuracy, completeness, merchantability, non-infringement of intellectual property, or fitness for any particular purpose. In no event will the author be liable for any damages whatsoever (whether direct, indirect, special, incidental, or consequential, including, without limitation, damages for loss of profits, business interruption, or loss of information) arising out of the use of or inability to use the information provided in this document, even if the author has been advised of the possibility of such damages.

Simon Southwell (simon@anita-simulators.org.uk)

Cambridge, UK, September 2021



Introduction

This document details the design and use of the open-source `rv32i_cpu_core` RISC-V softcore. It is a Verilog based design targeted at FPGAs.

It is meant as an exercise in implementing a modern RISC based processor with the aim to demonstrate the operations of such a device, and some of the design problems encountered in constructing such a component. The code itself is also designed to be accessible and expandable, as an education tool and a starting point for customisation. As such, this document, as well as usage information, will give some details of the internal design of the source code sufficient to navigate and modify the code, as desired. At this time it's fixed at implementing the RV32I features, with the extensions being a possible future expansion.

The implementation trades off between performance, resource usage and clarity of design. It was designed to be an efficient implementation running at a usefully high clock frequency, with a modest FPFA resource requirement—but where clarity of operation would be unnecessarily obfuscated, a more practical route is taken. The design can be used without the need to understand the code, but if one desires to delve deeper and make modifications, then the intent is that the code is as easy to understand as possible, consistent with implementing the desired features, and design details are well documented in this manual, along with some of the design decisions that were made along the way.

This being said, an initial specification, when targeting the chosen example platform (a DE10-nano with a Cyclone V 5CSEBA6U23I7 FPGA) is for a 100MHz clock, less than 1000ALMs (~2600 LEs) and a pipelined Harvard architecture.

Along with the implementation is an example simulation test bench, targeting the ModelSim simulator from Mentor Graphics, and the synthesis setup to target the DE10-nano, using the Intel/Altera Quartus Prime tools. These two environments are documented within this manual as well.

***rv32i_cpu_core* Supported Features**

The overall features of this RV32I initial delivery are listed below

- All RV32I instructions implemented
 - Configurable for RV32E
 - Single HART
- Separate instruction and data memory interfaces (Harvard architecture)
- 5 deep pipeline architecture
- 1 cycle operations for all instructions except branch, jump and load
 - Regfile update bypass feedback employed
- Branch instructions take 1 cycle when not branching, 4 when branching
 - 'never take' branch prediction policy employed, with pipeline cancellation on branch
 - Jump instruction takes 4 cycles
- Load instructions take a minimum of 3 cycles, plus any additional wait states
- Register file configurable between register or RAM based
 - defaults to RAM based, using $2 \times M10K$ RAM blocks
 - Register based costs approximately 700ALMs (~1900 LEs)
- Example simulation test bench provided
 - Targets ModelSim
- Example FPGA target platform using the terasIC [DE10-nano](#) development board (employing the Intel [Cyclone V 5CSEBA6U23I7](#) FPGA).
 - Targeting 100MHz clock operation

- < 1000 ALMs (~2600 LEs) when also employing Zicsr and RV32M extensions (RV32I implementation currently around 650 ALMs, ~1800 LEs).

The initial delivery instantiate the design in a core component, suitable for adding into the FPGA environment as a “Platform Designer” component. This core also instantiates some local FPGA RAM for the instruction memory (imem) and data memory (dmem). A future enhancement is to add a memory management unit for use of the, up to 1GB DRAM available on the DE10-nano.

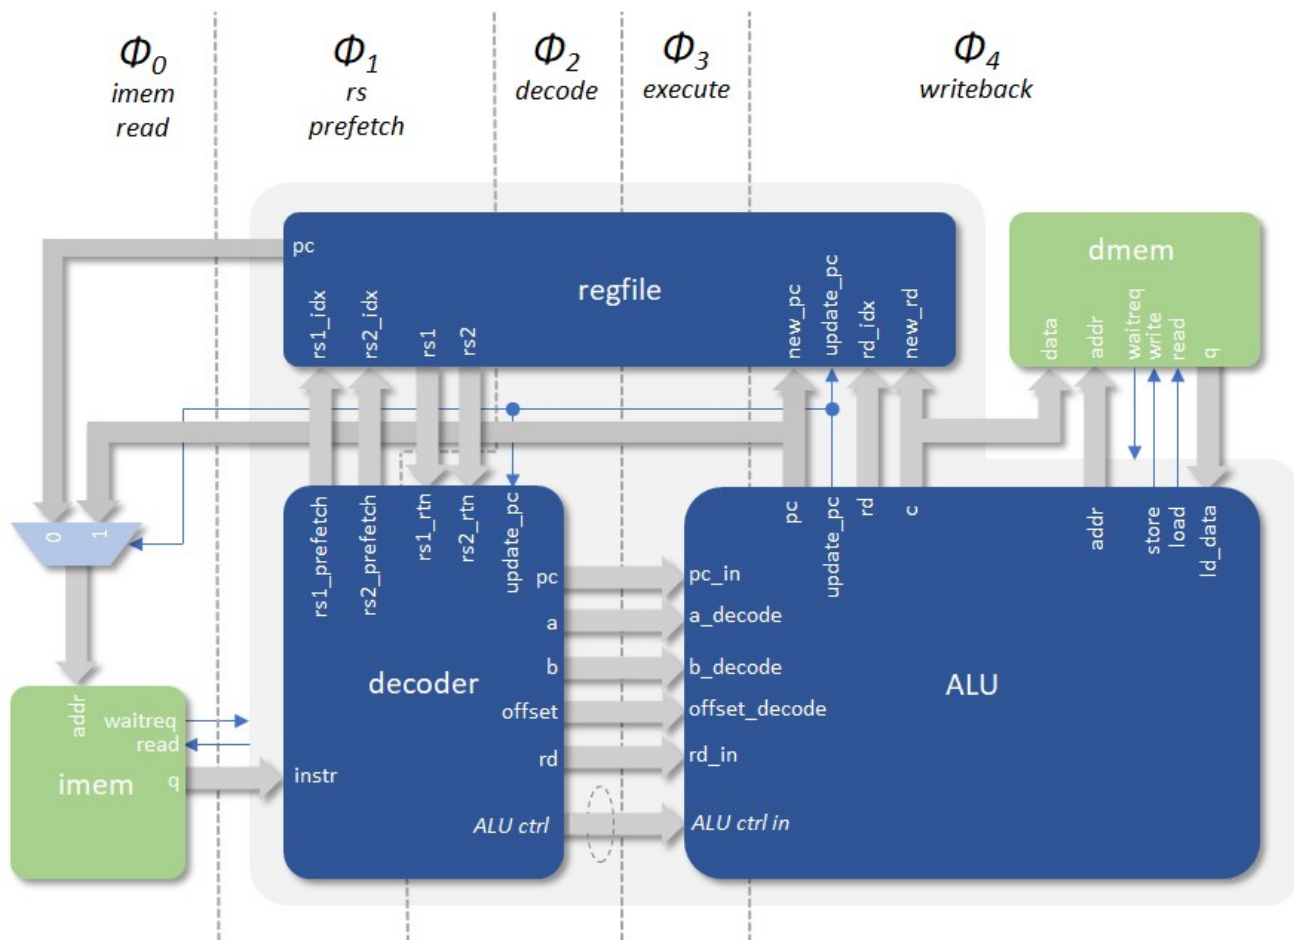
RV32I Core Implementation

The base core implements the RISC-V RV32I specification (see [1]). The idea is to implement a 'classic' RISC processor architecture, with multi-stage pipelining, to allow all instructions that have inputs and outputs only within the core, and do not alter the flow of the running program, to execute in a single cycle. That is to say, each instructions takes as many cycles as the pipeline is deep, but that the pipeline stages each have a partially executed instruction at that stage, so that the rate of processing is 1 cycle per instruction—with just an initial latency to fill it.

The instructions that do not have single cycle operations are jumps, branches (if taken), and memory load instructions. Memory store instructions will also have wait states when a memory sub-system is added but, for now, the example design that instantiates the core, uses FPGA memory blocks, which can be written in a single cycle.

Top Level

The implementation of the `rv32i_cpu_core` consists of three modules, which implement a 5-stage pipeline. The diagram below shows a simplified view of these components, which are instantiated within a top level.



The source code for the CPU can be found in the repository at `HDL/src`, and consists of the following files

- `rv32i_cpu_core.v` : Top level CPU core
- `rv32i_regfile.v`: Register file, containing the CPUs registers and program counter (the first phase of instruction reads)
- `rv32i_decode.v` : Instruction decoder implementing the next two phases of the pipe, source register pre-fetch and instruction decode
- `rv32i_alu.v` : Arithmetic logic unit executing the instruction and register write-back (that last two phases)

Pipeline Overview

The 5-stage pipeline allows for the efficient implementation of instructions, by breaking up the individual calculations required for the instructions into smaller pieces. It would be possible to execute an instruction in a single clock cycle, but that clock cycle would have to be long enough for all the calculations to be completed before allowing the next instruction to proceed. The aim of the pipeline is to divide the calculations into smaller 'chunks', with the output of one stage, feeding its results to the next stage. Since each stage's calculations are smaller than the whole, the clock cycle speed may be increased, limited by the slowest of the stages. In an ideal design, the stages would all take the same time, and the pipeline would be balanced, and run as fast as is possible.

Register Hazards

Having a pipeline architecture cause its own problems, however. If an instruction's input relies on a result of a neighbouring instructions before it, but it takes several cycles before that result is available (in a register, say), then that instruction may get stale data if it does not wait for the result of the instruction in front of it to be available. This gets worse, the deeper the pipeline. Several options are available. Conceptually, the simplest system might be that the issuing of instructions is halted whilst an instruction is in the pipeline, and is only issued when the writeback phase is reached for the instruction in front. This, though, gets back to the single cycle operation, as only one instruction is active at a time. Overheads would probably make this worse than single cycle operations. A better approach might be to know which registers (or memory locations, possibly) are to be updated by the instructions in the pipeline, and issue a new instructions only if there are no hazards—that is, its source registers do not match any of the destination registers of instructions in the pipe. This would require keeping and monitoring destination registers (and, possibly, store addresses), at each stage.

The approach take in this design it to allow instructions to proceed regardless. The values of the registers indexed by the instruction's RS1 and RS2 indexes (as appropriate), are accessed very early on, with the possibility of retrieving stale value. At each subsequent stage, except writeback, if a destination register is being written back, the stage will swap its source value with the writeback value, if its index matches that being written (and is not `x0`, which is never written). Since the ALU's execute phase immediately precedes the writeback phase, there is opportunity to update the values at all the phases up to that point. The details of destination feedback will be discussed on the sections for each block.

Branch Hazards

Another issue with a pipeline that allows instructions to always be issued is the problem of jumps and branches. If a jump is issued, or a branch that will be taken, the PC is updated from a linear increment to some arbitrary address to start fetching instructions from that location. When the jump is issued, though, it will be several cycles before the PC is updated at the writeback phase. In the meantime, the pipeline is fetching instructions beyond the jump, as if the jump has not happened. Those are not meant to be executed, and can alter state, with dire consequences.

In this design, the instructions proceed until an update to the PC is flagged at writeback. At this point, all the stages in the pipeline are 'cancelled'—effectively they are turned to nops (no-operation), and the pipeline continues. Since state is only updated at the writeback phase, by

forcing the stages to nop, none of the erroneously fetched instructions have any effect, even though the pipeline proceeds. What this does do, though, is introduce cycles in the pipeline that do nothing. This is why, in the sections above, it has been stated that branches take 1 cycle if the branch is not taken but 4 cycles if taken. This is a function of the pipeline depth, and it will be seen in other implementations too. Since the jump instructions are effectively a branch with an 'always taken' status, they are handled no differently, and the pipeline is 'stuffed' with the next instructions. The new PC value will not be known until the writeback stage, so early decode of jumps to stop issuing instructions until the new PC is available would just introduce additional logic to handle these. This might have an advantage in a power-conscious design, when one would not want to clock stages that do nothing, but this design does not attempt to be that sophisticated.

Stalling Mechanism

Another hazard with a pipeline architecture is for instructions that can cause "wait-state" or stalls. This boils down accesses to memory (and memory mapped I/O). In a general system, a memory management unit (MMU) might be employed which gives access to the DRAM (say) for a large storage space. Accessing that space can take tens of cycles (or more if I/O). A cache can be employed, with local fast access memory, with, perhaps, single cycle accesses. But, eventually, even the cache will need to be updated from main memory, and the access by the CPU will have to stall. This applies to both reading instructions from memory as well as load and store instructions.

The current design treats instruction reads and load/store instructions differently. The example design using the rv32i_cpu_core, has local FPGA memory, with single cycles writes, and a delayed cycle read. The design, however, has the ability to handle arbitrary wait-states for both instruction reads and memory access instructions.

imem reads

For instruction reads the, design policy is to issue nop instructions for all cycles where instruction read data is not valid. This takes care of situations where the memory sub-system stalls on instruction fetches, and in the initial latency of fetching the first instructions.

As mentioned before, this is not perhaps the best policy from a power consumption point of view (if a design has features to support this), but, for this design, it is the simplest solution.

dmem loads

For memory accesses, the example design, with local memory will not stall on store instructions. On load instructions, however, the reading from memory takes a single cycle. In addition, the address to be read is not available until the execution phase output, as it is calculated during this phase. So wait states must be introduced to allow issue the read, wait for the return value, and then writeback to the destination register.

Conceptually, the easiest thing to do might be to stop issuing new instructions after a load (or stuff with nops), until the load writeback is reached. Decoding for a load instruction takes a couple of phases, and so a latency on stopping the PC increment exists, and so some sort of wind-back might be needed to pick up at the correct address when the load is complete. The problem with this approach is that it empties the pipeline behind the load instruction, adding additional cycles over-and-above the wait states.

In this design, the approach is to allow the pipeline to proceed behind the load instruction, until the execute phase, when the pipe is stalled. The writeback phase will hold with the memory read accesses for as long as necessary to get the returned data, and route it to the destination register. At this point the pipeline is allowed to proceed, and no added latency was introduced by stuffing the pipe with nop instructions.

The spreadsheet fragment below is an attempt to show what happens for a load access stall (and was used for the signalling design)

		additional waitstates									
	imem read (pc)	$\Phi 0$	ld (addr+0)	addr+1	addr+2	addr+3	addr+3	...	addr+3	addr+4	addr+5
	rs prefetch (instr)	$\Phi 1$	-	ld	l1	l2	l2	...	l2	l3	l4
	decode(instr_reg)	$\Phi 2$	-	-	ld	i1	l2	...	l2	l2	l3
	execute (a, b)	$\Phi 3$	-	-	-	ld	i1	...	l1	l1	l2
	writeback(c, rd)	$\Phi 4$	-	-	-	-	ld	...	ld	ld	l1
dmem signals	read		0	0	0	0	1	...	1	0	-
	wait		0	0	0	0	1	...	0	0	-
ctrl signals	hold PC		0	0	0	1	1	1	0	0	-
	hold instr_reg, a+b, c, rd		0	0	0	0	1	1	1	0	-
	clear load		0	0	0	0	0	0	1	0	-
			Normal (unstalled) pipeline state								
			dmem wait state								
			hold state								
			mem access data valid								
			Set low by 'clear load' (as held load state would otherwise set it, starting a new access)								

The spreadsheet shows how the PC and phases proceed until stalled. A stall is 'initiated' when a dmem access is activated (read signal = 1), but the wait signal from the memory is also high. This is shown in read, and this may be the state for a number of cycles. Now it is known that a load instruction will have at least one wait state, and so early warning on stalling the PC can be affected when the load reaches the execution phase. This is needed so as not to push instructions into the pipeline with no space left, and hence lose them. The 'hold PC' column shows this signal, asserted before the others.

The other stages in the pipeline need to stall on the first memory access wait state, and continue to stall until valid data is returned (when 'wait' is de-asserted—green). This is shown in the row below 'hold PC'. Finally, the next cycle after the memory read valid data is where the writeback to a register is performed. Since the load instruction is still active, the dmem 'read' must be cleared in this phase (which would otherwise do because of the still active load instruction), so a new access is not started.

After this, the pipeline continues, with the instructions behind the load already to go. Thus only two additional cycles are added.

Parameters and Configuration

In general, the logic will work as intended without the need for configuration. However, the CPU core (rv32i_cpu_core), and the top level wrapper (core) have some parameters to alter their configuration, and these are documented below.

Parameters for the rv32i_cpu_core

The following lists the parameters available for the CPU core module, which are also routed to the core top level to make available to test benches and synthesis instantiations.

- RV32I_TRAP_VECTOR
- RV32I_RESET_VECTOR
- RV32I_REGFILE_USE_MEM
- RV32I_LOG2_REGFILE_ENTRIES (5 for RV32I, 4 for RV32E)
- RV32I_ENABLE_ECALL

The initial RV32I softcore implementation does not have CSR registers, as per the Zicsr extension specification, and so does not have an mvtc register defining the address for traps. The RV32I_TRAP_VECTOR parameter defines this address and has a default value of 0x00000004. Since there is no mvtc register, this cannot be changed programmatically. When the Zicsr extensions are added, this will become the default value of the mvtc register.

In a similar manner, the reset vector is defined using the RV32I_RESET_VECTOR parameter, and defaults to 0x00000000. Unlike the trap vector, there is no CSR register to change this, and the specifications leave this as an implementation choice. The two default values for the vectors, for this design, have been chosen to match the riscv-tests expectations.

The RV32I_REGFILE_USE_MEM parameter is used to select between a register based register file array when 0, or a RAM based register file array when 1. The default is 1 for RAM based. The choice is based on whether register or RAM is a limiting factor in a design using the core. In the current design, with a register based regfile, the FPGA resources used for the array is around that used for the rest of the core, doubling the requirements. If RAM is chosen, the two M-RAM blocks are used. For the example target platform this is two M10K blocks, with only 1024 bits of the 10240 bits used in each block. There are a few hundred blocks in the target FPGA, however, so the percentage usage is very low. If the core were targeted at an ASIC development, though, RAM may be more of an overhead than gates, so both implementations are available.

The RV32I_LOG2_REGFILE_ENTRIES parameter defines the 'address' width for the register file register array, which, thus, defines the number of entries. The default value is 5, giving the RV32I specification of 32 registers (x0 to x31). If an RV32E specification is required, this may be altered to 4, for 16 registers (x0 to x15). When selecting a RAM based register file array, switching to an RV32E specification will save no RAM bits, and so is less useful.

The RV32I_ENABLE_ECALL parameter is purely for test purposes. It is there to allow the riscv-tests code to be run without the need for Zicsr extensions. These tests, as part of initialization set up some CSR registers and make tests on others. Normally, tests are made and will skip part of the initialisation, including executing an ecall instruction. With the parameter set to 0 (it defaults to 1), the ecall instruction is treated as a nop. With the implementation's default values, and setting this parameter, the test code drops through to the main test code with sufficient initialisation to execute without issue. The intent is to remove this parameter once Zicsr extensions are added.

Additional Parameters for the top level Core

As well as the parameters for the CPU core, which are available at the top level core, some additional parameters are added, as listed below.

- RV32I_IMEM_ADDR_WIDTH
- RV32I_DMEM_ADDR_WIDTH
- RV32I_IMEM_INIT_FILE
- RV32I_DMEM_INIT_FILE
- RV32I_IMEM_SHADOW_WR

The RV32I_IMEM_ADDR_WIDTH and RV32I_DMEM_ADDR_WIDTH parameters specify the address widths of the internal instruction and data memories, and hence their size. By default they are 12 bits, for 4096 x 32 giving 16Kbytes for each RAM.

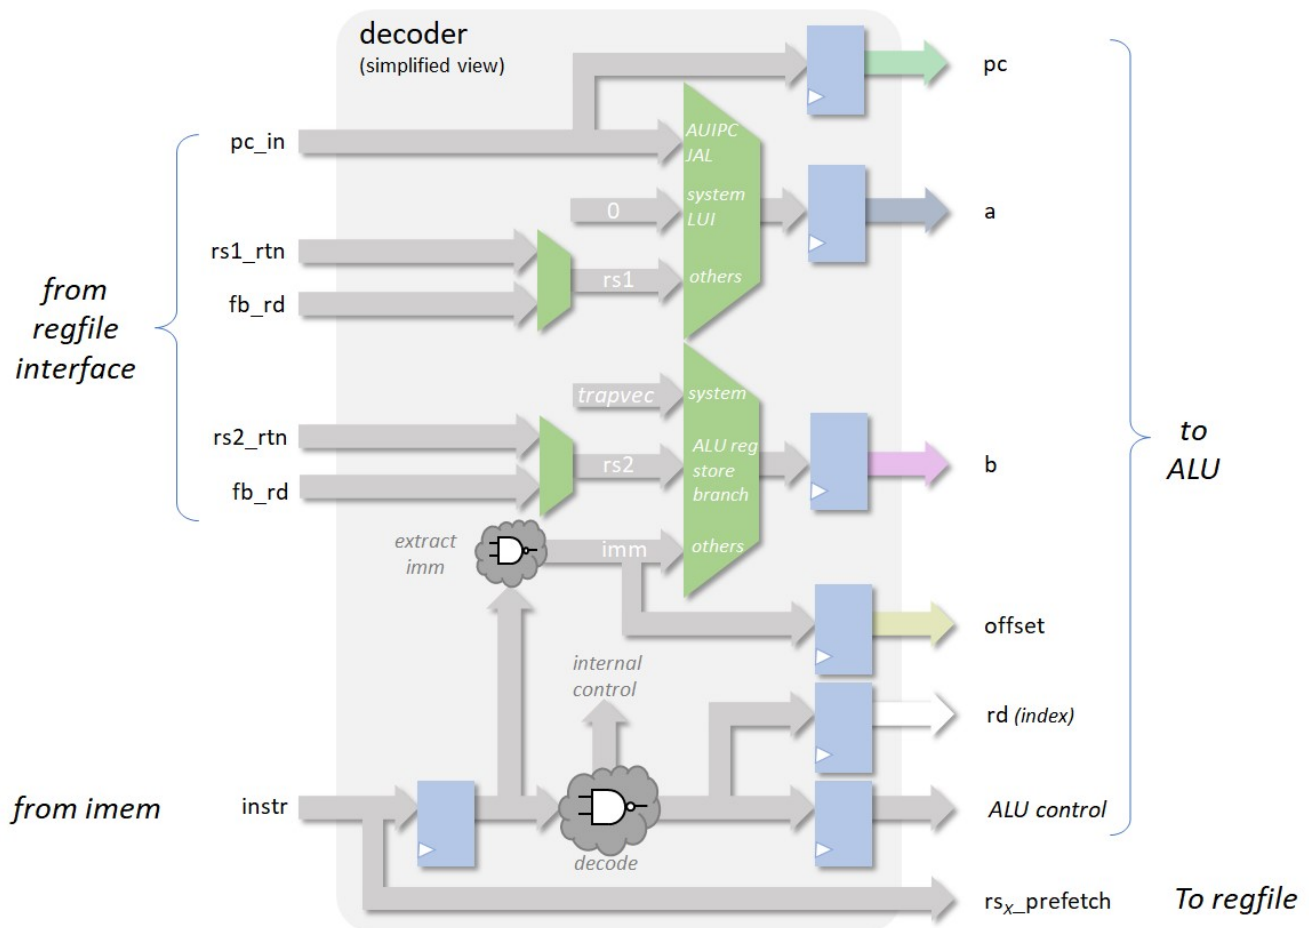
Memory initialisation files can be specified for the RAMs with the RV32I_IMEM_INIT_FILE and RV32I_DMEM_INIT_FILE parameters. The format for these are specific to the targeted example platform, and can either be in intel hex format (.hex), or in a 'memory initialisation file' (.mif)

format. It is expected that, for different targets, the RAM modules will be replaced with the appropriate substitutes. The parameters can still be routed to these new blocks and specify an initialisation file of the appropriate type, if required. The default value for these parameters is "UNUSED", indicating no initialisation file is specified for the RAM.

The RV32I_IMEM_SHADOW_WR parameter is meant for test purposes only. By default it is 0, with the logic's behaviour as normal. When set to 1, logic is enabled where all writes to the data memory, will also be written to the instruction memory. This was added specifically to allow the riscv-tests fence_i.s test to write, where stores to code locations are made for self-modifying purposes to test the fence.i instruction. Since IMEM and DMEM are separate memories, writes to code locations will be routed to DMEM. By enabling this feature, these writes are shadowed in the IMEM as well. It is expected that in an implementation this parameter is left at its default value. The example synthesis environment has modification of this parameter disabled.

Decoder

The decoder block, implemented in rv32i_decode.v, is responsible for processing raw instructions values and splitting out the instruction type and their various relevant fields of the instruction for sending to the arithmetic logic unit (ALU), as well as instigating source register prefetching. Thus it implements the pipeline stages 1 (rs prefetch) and 2 (decode). The diagram below shows a simplified view of the decoder module.



Instructions arrive from IMEM on the `instr` input, and the logic immediately routes the RS1 and RS2 fields to the register file block to start retrieving the values. This is regardless if what

type the instruction will decode to be. The RISC-V specification has been constructed in such a way as to position the RS indexes in the same place for each instruction that specifies a source register. By routing the bits in phase 1 to the register file block for all instructions, without further decode, the values become available for phase 2, if required. Otherwise the returned values are ignored. For a power sensitive design, additional decode in phase 1 might allow only fetching of data when relevant. Since the instr input is from memory, then it may arrive late in the cycle, and so needs to be only lightly processed where possible.

The instruction input is registered for the phase 1 processing. There are three main sections to the decode. Firstly, all possible opcode/function, register indexing and immediate value fields are extracted from the instruction. This is just wiring, and so is a convenience only.

The second stage is to use the extracted opcode and function fields to generate bits flagging the different instruction types, such as arithmetic, jump, memory accesses etc. With the operation bits, the various immediate fields are multiplexed to a single immediate value, based on the decode instruction type.

The final stage is to generate the outputs for routing to the ALU. The main outputs are an A and B value for inputs to the main ALU calculations. The values set on these are dependent on the instructions type. A can be the RS1 value, the PC or 0. The B value can be the RS2 value, the trap vector value or the selected immediate value. With appropriate A and B values the ALU calculations can take these common inputs and generate its output. It will still need to know the operation to perform on these value, and so the operation bit values are sent to the ALU (shown as ALU control on the diagram).

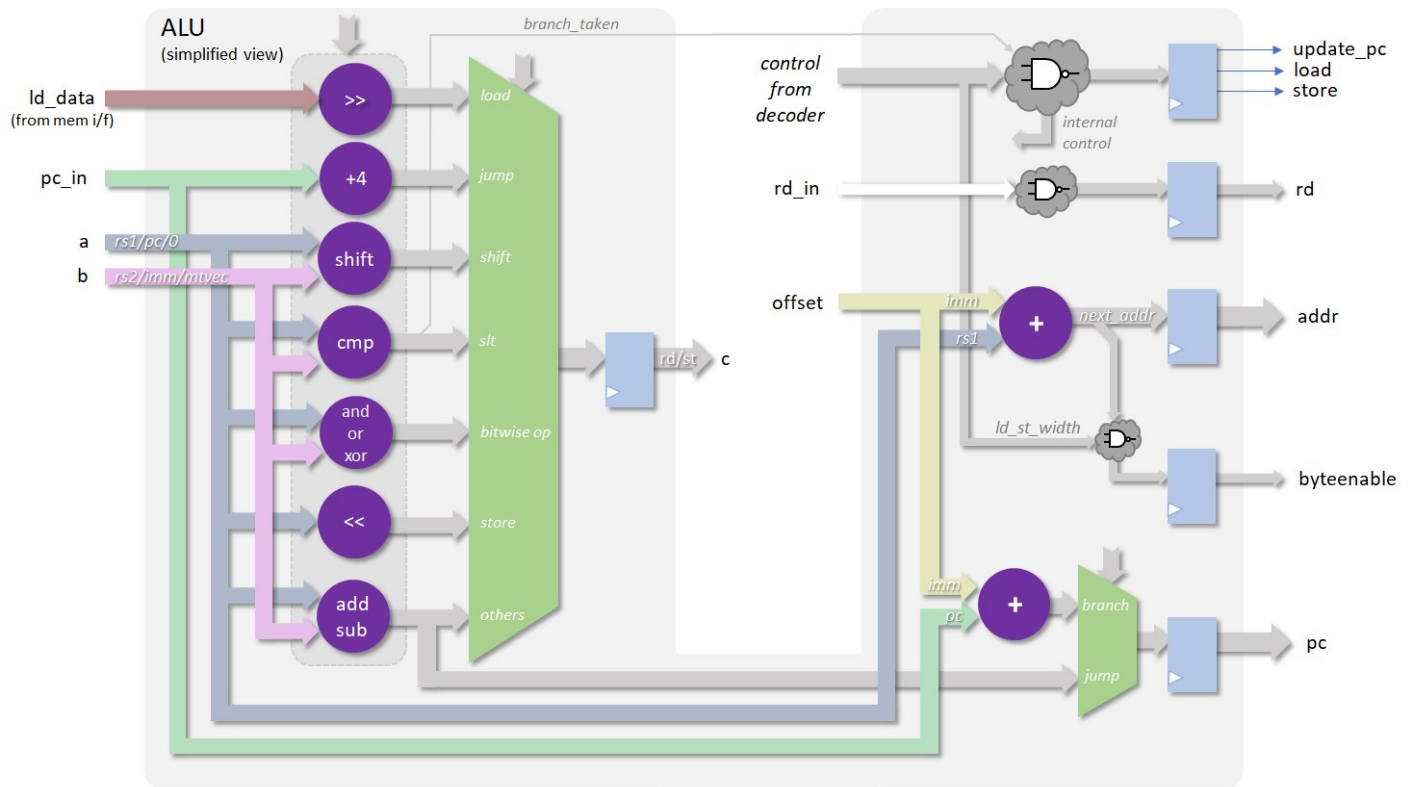
In addition to the A and B values, the source register indexes for the A and B values are sent to the ALU as well for instructions that use them (not shown in the diagram). This is purely for the pipeline destination (RD) feedback mechanism to allow for updating stale RS fetches (see pipeline section above). The decode block itself monitors for matching RS and RD indexes, and substitutes the register file returned values for the RD value update when a match occurs. This is not shown in the diagram. The RD index is also sent to the ALU to allow the ALU result to be routed to the appropriate destination register.

In addition to the A and B values, some instruction require calculations on the PC. For example, branch instructions must do a comparison of two registers, and A and B are used for these values, but a new PC is required if the branch is taken and a calculation is required on the PC value an offset value. The PC and the extracted offset are also routed from the decoder to the ALU.

Arithmetic Logic Unit

The ALU is where all the calculations for an instruction are executed, and where the results are written back to registers or memory, where applicable. New address values are also written back to the PC when necessary. Thus it implements phases 3 and 4 of the pipeline.

The diagram below shows a simplified view of the ALU logic.



At the heart of the design is the main execution calculations, when take the A and B values from the decoder to provide a result in a C output. For the most part, A and B are fed to the various possible operation types, and the appropriate result mux'ed to the C register, selected from the instruction type bit controls from the decoder. The categories of instruction in this mux are:

- load
- store
- jump
- shift
- comparison (not branches)
- bitwise operations
- arithmetic (add and subtract)

Of these categories all use the A and B inputs, except the following, The load instruction takes the returned data and aligns the data dependent on the data width and address. Similarly, the store instructions have the source data aligned to the appropriate byte lanes, ready for writing to memory. Jump instructions take the PC value for the instruction and adds 4. This is for the link address that will then be written to the indexed destination register, The actual new PC value is calculated separately and is discussed below. All the other types of instruction take the A and B values, as selected by the decoder, whether from a source register, and immediate value, or whatever. Not shown in the diagram is the writeback feedback mechanism where the A or B values from the decoder are overridden by the C output if their index matches the destination register index (and not 0). This is the phase 4 stage's 'stale rs data' mechanism,

As mentioned before, some of the instructions require additional calculations. For loads and stores a memory address is calculated. This takes the an address base value on the A input and an offset value sent from the decoder and extracted from the instruction's immediate value. Normally an immediate value appears on the B input, but the main calculation logic is

busy with either aligning memory read data (loads) or aligning write data (stores). An `addr` output sends the address to memory, and with byte enables calculated from the bottom bits.

For jumps and branches, a new PC value must be calculated, and this comes from the offset mentioned above and the PC value input to generate a new PC value.

All the `c`, `addr` and `pc` calculations proceed regardless of instruction type being executed. Whether they affect state, or not, depends on the control outputs. A writeback to a destination register only happens if the `rd` index output is non-zero. Also, an `update_pc` output flags to write the `pc` output to the PC register. A load and store output control initiating accesses to memory. All of these are functions of the instruction type control inputs.

Stalling

In the pipeline section, stalling was detailed from an overall perspective. The calculations on when to stall each stage of the pipeline are done at the top level. For the ALU, a single stall input controls stalling of the phase 3 and 4 stages. When stall is asserted the ALU will hold key state with its previous value. These include the following outputs

- `addr`
- `load`
- `store`
- `rd`
- `update_rd`
- `pc`
- `update_pc`

In addition, some internal state, `ld_width`, is also held which the byte enable outputs are calculated from.

Register File

The register file module is where all the main CPU state is implemented---the main integer registers `x0` to `x31`, and the PC program counter. At its simplest, it is an array of 32 bit registers that can be written and read. As it holds the program counter, it implements phase 0 of the pipeline.

Registers versus RAM

In this implementation, the registers can be configured to be implemented as an array of logic registers or as memory blocks under the control of a parameter, `RV32I_REGFILE_USE_MEM`.

When registers are chosen, an array of 32 bit registers is instantiated, which allows then to be written, selected with an `rd_idx` input and a new `rd` value to be written. Two values need to be read simultaneously, for RS1 and RS2 values, and `rs1_idx` and `rs2_idx` index the registers to be read, with data returned on `rs1` and `rs2` outputs in the following cycle. Note that all three ports can be active simultaneously.

If memory blocks are used, to save on register usage, some problems arise. A two port memory block can have, at most two active accesses, where three are required. Secondly, write-through (where a location being read that's also being written sees the new value) is not possible without additional logic. These problems need to be solved if RAM blocks are to be used.

The multiple accesses issue is solved by instantiating two RAM blocks, which both get written simultaneously when a new RD value is being written, thus they always contain the same data.

When two reads are required, one is directed to the first RAM, the second to the other RAM, giving simultaneous read access. This, however, does not solve the write through problem.

Pipeline Feedback

When writing to the RAM, the inputs are latched, and then the value is written in the next cycle. If reading in the same cycle as the write, these inputs will also be latched, and data read in the next cycle, but this will be the old data, as the state is updating in the same cycle. So, the regfile module has a short queue of recently updated values to bypass the RAM in these cases.

If a read access matches the input write location, the value is simply passed back as the read value, just as for pipeline feedback in other stages. If the read index matches the previous write location, then a stored write value is routed back as the read value, to compensate for the cycles needed to update memory. In this manner, the correct value is returned, however close read and write accesses to the same location are.

This mechanism is employed whether registers or RAMs are used so as to keep the logic the same in both cases, thus simplifying it. It also allows the register implementation to register its returned value outputs, relieving timing for the decoder.

PC and Stalling

For the regfile module, stalling is just a matter of holding the PC state. Two PC values are kept in the register file block. The pc output is that used to fetch the instructions from IMEM. An additional value, `last_pc`, is the value of the PC in the last cycle (normally PC - 4). This value is fed to the decoder as the PC value of the instruction sent, since there is a delay in instruction fetches, and is

Memory Interfaces

The two memory interfaces, for the instructions (IMEM) and load/store data (DMEM) are modelled on the Avalon memory mapped master bus specification from Intel/Altera. This is because the example target platform uses the Intel/Altera Cyclone V FPGA, and this maps well to a solution based on these devices. Mapping to a more generic protocol, such as the open source Wishbone interface should be quite simple, with the main difference being a 'cycle' signal to indicate a bus access, and a write enable indicating a write (and not read), instead of separate read and write strobes.

The signalling is not fixed cycle but has a wait request input back to the master buses of the CPU core, to introduce wait states as discussed before for loads and instruction reads.

The current example design uses local memories, but the longer term aim is to have a memory sub-system with caches and access to main DRAM memory. The use of wait request inputs allows for this enhancement, with arbitrary wait states that may occur, without alteration of the CPU core logic.

Verification

The verification strategy of the design follows that of the instruction set simulator. The [riscv-tests](#) supplied by RISC-V International, which are categorised into tests for the various base specifications and extensions. They are self-checking assembly code, and thus require a minimum amount of functionality in order to run correctly. These tests, in turn, rely on the [riscv-test-env](#) repository.

Some ‘turn-on’ test code was written in order to get to a useful standard, but this was not self-checking and is superseded by the riscv-tests code, and so is not part of the repository.

An example simulation test environment is provided to run the riscv-tests, as well as a synthesis environment to build for the DE10-nano board, with the ability to run these same tests. The following sections describe these environments and how they are used.

Simulation

The CPU core simulation environment can be found in the repository in the folder HDL/test. This contains all the files to compile and run a simulation for the core.

Compiling and Running the Simulation

Prerequisites

A provided makefile allows for compilation and execution of the example simulation tests. The current support for simulation is only for ModelSim at this time, and has some prerequisites before being able to be run. The following tools are expected to be installed

- ModelSim Intel FPGA Starter Edition. Version used is 10.5b, bundled with Quartus Prime 18.1 Lite Edition. The MODEL_TECH environment variable must be set to point to the simulator binaries directory
- MinGW and Msys 1.0: For the tool chain and for the utility programs (e.g. make). The PATH environment variable is expected to include paths to these two tools’ executables

The versions mentioned above are those that have been used to verify the design, but different versions will likely work just as well. If MODEL_TECH environment variable is set correctly, and MinGW and Msys bin folders added to the PATH, then the makefile should just use whichever versions are installed.

Compiling test code

Before running a simulation a test file must be created to load into the instruction memory. By default the memory is looking for a memory initialisation file named test.mif. This will be generated from one of the tests in the [riscv-tests](#) repository. The riscv-tests repository is expected to be checked out in the same directory as this repository, as well as the accompanying [riscv-test-env](#) repository, which riscv-tests relies on.

A make file (makefile.test) is provided to compile a test to an executable and then convert this to a memory initialisation file (.mif). There is a program supplied for the ELF to .mif conversion, which is in the test folder and called elf2vmif.cpp. The make file takes care of building this if it hasn’t been built already.

The make file itself is executed with some arguments to override variables to select a test to be compiled and the sub-directory in riscv-tests/isa where it is located. For example, from a command window, and in the test folder, one can execute, say:

```
make -f makefile.test SUBDIR=rv32ui FNAME=addi.S
```

This will compile the `addi.S` test to a file in the test folder called `addi.exe`. It then converts this to a file `test.mif`. If a simulation is now run, the compiled code will be loaded into the instruction memory before the simulation is run.

Running a Simulation

The main make file (`makefile`) is used to compile and then run a simulation, either as a batch simulation or a GUI simulation. Typing `make help` displays the following usage message:

```
make help      Display this message
make           Build C/C++ code without running simulation
make compile   Build HDL code without running simulation
make run/sim   Build and run batch simulation
make rungui/gui Build and run GUI simulation
make runlog/log Build and run batch simulation with signal logging
make waves     Run wave view in free starter ModelSim (to view runlog signals)
make clean     clean previous build artefacts
```


Giving no arguments compiles the Verilog code. The `compile` does the same thing. The various execution modes are for batch, GUI and batch with logging. The batch run with logging assumes there is a `wave.do` file (generated from the waveform viewer) that it processes into a `batch.do` with signal logging commands. When run, a `vsim.wlf` file is produced by the simulator, just as it would when running in GUI mode, with data for the logged signals. The `make wave` command can then be used to display the logged data without running a new simulation. Note that the `batch.do` generation is fairly simple at this point, and if the `wave.do` is constructed with complex mappings the translation may not work. The `make clean` command deletes all the artefacts from previous builds to ensure the next compilation starts from a fresh state.

The `riscv-tests` exit with the `gp` register containing a test number from bit 1 upwards, and the bottom bit set. If the test number is 0, then the test has passed (if the bottom bit is set). Otherwise the test number is the number of a test that failed (it halts after the first failure). The test bench, at the end of the test inspects the value and prints a message for pass or fail, with some information about the state.

All of the `rv32ui` tests can be run as a regression test using the `run32i_tests.bat` batch file. This goes through all of the RV32I unprivileged instruction tests, compiling them, running a batch simulation, and extracting the PASS/FAIL messages in a test log (`test.log`), which it prints to the screen at the end for inspection for failures.

Test Bench

The test bench (`tb.v`) is a straight forward Verilog module, in a traditional format. The top level core is instantiated as the UUT, and a clock and reset are generated locally and sent to the core. Not much else needs to be connected to the core, and the main other inputs to the design are the parameters. The test bench module has parameters, some of which are those of the instantiated top level core, allowing them to be overridden when running a test. The diagram below shows the general structure of the test bench and instantiated core.



The core has an optionally instantiated test block for monitoring the `imem` read bus, and halting the core when certain conditions are met, as defined by the control register bits in the local registers. The test bench sets the `RV32I_INCL_TEST_BLOCK` parameter to 1, so that this block is present during simulations.

The core has an Avalon memory mapped slave bus for reads and writes to registers/memory from an external source. In normal usage, this is mainly for loading a program to instruction memory, and then taking the core out of reset to execute the program. In simulation, the program is loaded from an initialisation file to save wasting time writing to memory during simulation. Since it is instantiated, the core's test block is configured using other registers, to select halt conditions (on unimplemented instruction reads and/or specific address reads). When a halt condition occurs, the core will set a signal which can be read via the registers. The registers in the core are defined as shown below:

core		This block is the core registers				
Register Name	Word Addr Offset	Field	Bits	Type	Reset	Description
CSR_CORE_CONTROL	0x0		3			Core control register
		CLR_HALT	0	w0	0	Clear test halt status
		HALT_ON_ADDR	1	w	0	Enable/disable halt on specific address
		HALT_ON_UNIMP	2	w	0	Halt on unimplemented instruction
CSR_CORE_STATUS	0x1		2			Status of core
		HALTED	0	r	0	Halted status
		RESET	1	r	0	Reset status
CSR_CORE_HALT_ADDR	0x2		32	w	0x00000040	Halt address, if halt_on_addr active
CSR_CORE_GP	0x3		32	r	0	gp register value

Some probes into the core are made from the top level to extract signals for control of execution. In particular, the GP register (x3) and PC register, along with the `test_halt` signal from the test block. The test bench logic monitors these signals and determines when a halt condition occurs. It is when the test bench detects a halt condition that test status is inspected and a PASS/FAIL message is displayed.

Core Parameters

The core has various parameters which configure the logic. In the test bench, these are configured in such a way as to allow the riscv-tests to be run. The table below shows the parameters, their default values and the values set by the test bench:

Parameter	Default	Test Bench
CLK_FREQ_MHZ	100	100
RV32I_RESET_VECTOR	32'h00000000	32'h00000000
RV32I_TRAP_VECTOR	32'h00000004	32'h00000004
RV32I_LOG2_REGFILE_ENTRIES	5	5
RV32I_REGFILE_USE_MEM	1	1
RV32I_IMEM_ADDR_WIDTH	16	16
RV32I_DMEM_ADDR_WIDTH	16	16
RV32I_IMEM_INIT_FILE	"UNUSED"	"test.mif"
RV32I_DMEM_INIT_FILE	"UNUSED"	"UNUSED"
RV32I_ENABLE_ECALL	1	0
RV32I_IMEM_SHADOW_WR	0	1
RV32I_INCL_TEST_BLOCK	0	1

Results

The simulations were run for each of the riscv-tests under the rv32ui folder—that is, the unprivileged instructions for RV32I base specification. The table below shows the tests that have been run and their status.


sub-test folder	test	status
rv32ui	simple.S	Passed
	add.S	Passed
	addi.S	Passed
	and.S	Passed
	andi.S	Passed
	auipc.S	Passed
	beq.S	Passed
	bge.S	Passed
	bgeu.S	Passed
	blt.S	Passed
	bltu.S	Passed
	bne.S	Passed
	jal.S	Passed
	jalr.S	Passed
	lb.S	Passed
	lbu.S	Passed
	lh.S	Passed
	lhu.S	Passed
	lui.S	Passed
	lw.S	Passed

	or.S	Passed
	ori.S	Passed
	sb.S	Passed
	sh.S	Passed
	sll.S	Passed
	slli.S	Passed
	slt.S	Passed
	slti.S	Passed
	sltiu.S	Passed
	sltu.S	Passed
	sra.S	Passed
	srai.S	Passed
	srl.S	Passed
	srli.S	Passed
	sub.S	Passed
	sw.S	Passed
	xor.S	Passed
	xori.S	Passed

These are all the relevant tests for the current implementation. When the extension specifications are added, the additional tests will be run, as for the instruction set simulator.

Synthesis

An example target platform environment is provided, targeting the DE10-nano board (shown below). Files are provided for the top level core to be instantiated in a Platform Designer infrastructure (.qsys file) and the setup and constraints for a Cyclone V BA6 device synthesis, suitable for the chosen platform. These files can be found in the HDL/de10-nano/build folder.

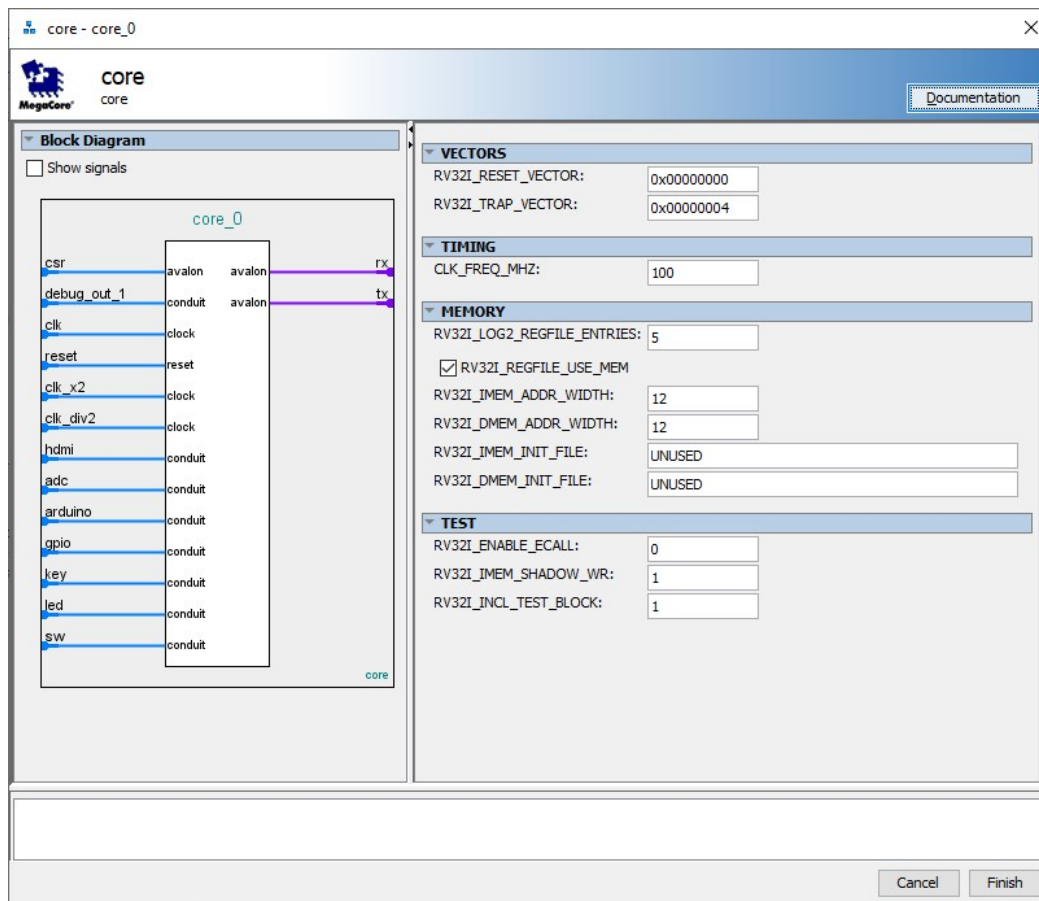


The environment assumes that the Quartus Prime tools have been installed on the host machine. The code has been tested on Quartus Prime Lite 18.1, but later versions should work just as well.

Core Component

The core logic is specific to the DE10-Nano target platform. Instantiates the `rv32i_cpu_core` and two Intel/Altera M-RAM memories. The source code is located in the HDL/de10-nano/src folder, along with a `top.v` file that is the synthesis top level for the target FPGA. The core has ports for all the possible interfaces for the DE10-Nano, but in this example most of these are not used.

The Quartus tools have been used to create a Platform Designer component from the core block that can be instantiated in the `top_system`, along with the Hard Processor System (HPS) and some other required peripherals. It is the `top_system` that is instantiated in the `top.v` code, and contains the core component. In the same folder as the `core.v` component file, the `core_hw.tcl` defines this component for use in the Platform Designer, listing the all the files required and other attributes. The parameters can be set in the platform designer by editing the instantiated core, and a windows appears as shown below with the main interfaces and the parameters for altering.



Command Line Build

The design can be synthesised using the Quartus GUI tools in a standard manner, but a make based command line build option is provided. In the build directory a makefile can be used to synthesise the design by simply typing make from a command window from within the build folder.

When this is done, the script will construct the Platform Designer top_system from the top_system.qsys file, and then proceed to go through the build steps of synthesis. When complete, the timing report will be scanned for violations, and any found printed out for inspection.

The generated reports and FPGA image files are all generated in an output_files folder. Both a top.sof and a top.rbf file are generated for use, as appropriate, to configure the FPGA.

Results

Timing

The design was constrained for the target clock frequency of 100MHz via the Platform Designer's instantiated PLL components, with outclk0 being this frequency and connected to the clk input of the core. No timing violations were reported and the estimated fmax frequency for the slowest corner was ~117.6MHz.

Resource Usage

The table below shows the achieved resource usage for the synthesis, showing both the actual number of ALMs used, plus an estimate of the logic element (LE) equivalent. Also the number of M10K RAM blocks is given.

module	ALMs (~LEs)	M10Ks
rv32i_cpu_core	6 (16)	0
rv32i_decode	133 (359)	0
rv32i_alu	480 (1296)	0
rv32i_regfile (RAM)	64 (173)	2
Totals	683 (1844)	2

Platform Testing

Testing on the example platform based around repeating the `riscv-tests` appropriate to the build—at this time, the `rv32ui` tests. In order to do this, the DE10-Nano FPGA needs to be configured with the image generated in synthesis (see Command Line Build section above). Some software is also required to load a program to the internal memory and then run a test.

The core component is configured to instantiate the test logic (see diagram in Test Bench section above), which monitors for configured halt conditions. It is controlled via local core registers, accessed over a memory mapped control and status register bus (CSR) and also makes available the halt status and the value of the GP registers (see register table in the Test Bench section above). The test software accesses these registers to load a program, bring the core out of reset and then monitor for a halt status. It then determines the exit status and prints a PASS or FAIL message. This mimics the running of the tests in the simulation test bench, and the core parameters are set in synthesis to be the same as for simulation, apart from the `RV32I_IMEM_INIT_FILE`, which is left at the default of "UNUSED", as the code will be loaded by test program. This is shown in diagram in the Core Component section above.

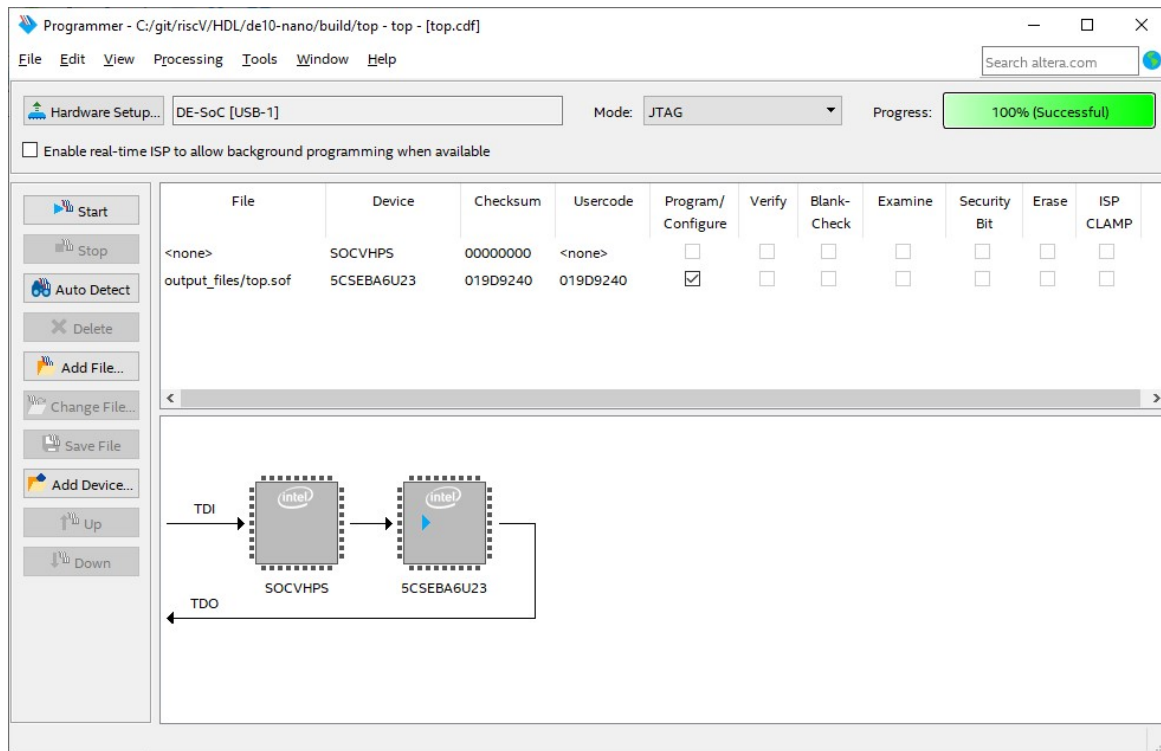
Loading the FPGA Image

This document cannot go through a tutorial on the use of the Quartus tools, and the procedures are specific to the Intel FPGAs and the DE10-Nano board. It is supposed that familiarity with these tools, or their equivalent if targeting a different platform, is sufficient without further discussion. A brief summary is documented here for the example platform to indicate the steps required.

The DE10-Nano, for these tests is running a console version of Ångström Linux. At time of writing, this image can be found [here](#). With this setup, there are two ways to load an FPGA image on the DE10-Nano; over USB-Blaster JTAG or copy an image to the FAT32 partition so that it is loaded at boot time. The later needs the FAT32 partition mounted (e.g. `/mnt/fat32`), and the generated RBF (`top.rbf`) copied to that folder and named `soc_system.rbf`. The advantage of this is that the FPGA will be configured with the image every time the board is powered on.

Here we will concentrate on loading over the USB-Blaster. From the Quartus Prime GUI, the Programmer window can be launched. If connected correctly, the tool will detect the connection. If the first time connected, then the Auto Detect button should be pressed to

inspect for the connected FPGA, but a top.cdf file is provided which already defines this chain.



In the diagram above is shown a programmer window with the correct chain configuration, and the top.sof file selected as the configuration image. Programming is then simply a matter of pressing the Start button. Programming this way means the image will be lost at power down, but it allows the image to be updated, without a reboot, when developing the logic.

Compiling the Test Program

The HDL\de10-nano\test folder contains the source files to compile a Linux ARM program to run riscv-tests/rv32ui tests on the DE10-nano platform.

A makefile exists to build the test program. On a windows host, it assumes that the following tool chain is installed at the given location:

```
c:\Tools\gcc-linaro-4.9.4-2017.01-i686-mingw32_arm-linux-gnueabi
```

If a toolchain is installed elsewhere, then the TOOLPATH variable should be modified, either on the command line, or the makefile updated. If no toolchain is available, a compiled version of the code is available in the folder (make.exe), but modifications to the program will not be possible. Alternatively, if the Linux running on the DE10-Nano has the GNU gcc toolchain installed, it may be compiled on platform. At the time of writing, the toolchain used in testing is available [here](#).

Uploading the Test Code

The main.exe executable should be copied to a suitable folder on the DE10-Nano. How this is done depends on the particular setup. For this development the DE10-Nano was connected to the local network (with a fixed MAC address—otherwise the allocated address must be discovered), and [MobaXterm](#) used to connect to the board as root. The diagram below shows a MobaXterm session and a single test run:

The screenshot shows the MobaXterm 20.3 interface. On the left, there is a sidebar with a file explorer showing the directory `/home/root/riscv/` containing files `main.exe` (17 KB), `run_tests.sh` (1 KB), `test.exe` (7 KB), and `test.log` (9 KB). The main terminal window displays the following output:

```
• MobaXterm 20.3 •
(SSH client, X-server and networking tools)

> SSH session to root@192.168.0.30
• SSH compression : ✗ (disabled or not supported by server)
• SSH-browser      : ✓
• X11-forwarding   : ✗ (disabled or not supported by server)
• DISPLAY          : 192.168.0.8:0.0

> For more info, ctrl+click on help or visit our website

Last login: Wed Sep 29 09:45:48 2021 from rioja.local
root@socfpga:~# cd riscv/
root@socfpga:~/riscv# ./main.exe

*****
*   Wyvern Semiconductors   *
* rv32i_cpu_core (ARM Cortex-A9) *
*   Copyright (c) 2021     *
*****

Test exit code = 0 : PASS

main() finishing with status 0

root@socfpga:~/riscv#
```

At the bottom of the terminal window, there is a status bar that reads: "UNREGISTERED VERSION - Please support MobaXterm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net>".

When run on the platform, `main.exe` program expects a RISC-V executable file called `test.exe`, which it will load to the internal memory and then execute. When the internal test block has halted, it will extract the state and print a PASS or FAIL message, as shown above.

Running the Tests

A script (`run_tests.sh`) is also provided to run the `riscv-tests/rv32ui` tests as a regression. A `test/` folder should be located in the same folder as the `main.exe` and the script, and contain all the compiled `rv32ui` tests. The script will run each executable in turn, and log the output in `test.log`. When complete all the PASS/FAIL messages are dumped to the screen for inspection.

The current version of the core has had all the tests run that were run in simulation (see table in the Results sub-section of the Simulation section), and all pass, verifying the logic on the example platform.

References

- [1] *The RISC-V Instruction Set Manual, Volume I: User-Level ISA, Document Version 20191213*, Editors Andrew Waterman and Krste Asanović, RISC-V Foundation, December 2019.
- [2] *The RISC-V Instruction Set Manual, Volume II: Privileged Architecture, Document Version 20190608-Priv-MSU-Ratified*, Editors Andrew Waterman and Krste Asanović, RISC-V Foundation, June 2019